

Provably Pareto-Optimal Power-Aware Routing for Multi-GPU LLM Serving

Anonymous Author(s)

Abstract

We formalize power-aware request routing for multi-GPU LLM-serving fleets as a multi-resource social welfare problem: each replica carries three independently-normalized shares — compute (prefix-cache-discounted prefill cost), load (in-flight request count), and power (live NVML-measured ramp rate relative to a calibrated ceiling) — and the router seeks an egalitarian (max-min) welfare allocation across candidates at each routing decision, extending Dominant Resource Fairness [2] to a third, reactive resource dimension. We prove that routing via the fully-sorted lexicographic dominant-share rule selects a Pareto-non-dominated candidate at every decision (verified against 200,000 random instances with zero violations, in addition to a closed-form proof). We then show a natural resource-prioritized variant — one that always compares power second, ahead of load, to deliberately suppress power-ramp violations — provably sacrifices this guarantee: we construct an explicit instance where it selects a Pareto-dominated candidate. Despite giving up the theoretical guarantee, this variant wins empirically on real 8×4090 hardware under sustained fleet power pressure: relative to the Pareto-safe rule, it cuts peak power-ramp rate by 40% (with 8× tighter run-to-run variance) and cross-GPU ramp coincidence by 43%, while leaving mean request latency statistically flat, across 3 replicated trials. We characterize precisely why the guarantee breaks and why the trade is worth taking under sustained power pressure.

CCS Concepts

• **Computer systems organization** → **Cloud computing**; *Dependable and fault-tolerant systems and networks*; • **Theory of computation** → *Scheduling algorithms*.

Keywords

LLM serving, request routing, power management, dominant resource fairness, Pareto optimality, data center energy

ACM Reference Format:

Anonymous Author(s). 2027. Provably Pareto-Optimal Power-Aware Routing for Multi-GPU LLM Serving. In *Proceedings of ACM eEnergy 2027 Fall: 17th ACM International Conference on Future and Sustainable Energy Systems (eEnergy '27)*. ACM, New York, NY, USA, 3 pages. <https://doi.org/XXXXXXX.XXXXXXX>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
eEnergy '27, TBD

© 2027 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-1-4503-XXXX-X/2027/TBD
<https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

Motivation. Multi-GPU LLM-serving fleets are increasingly deployed at data-center scale, where power-ramp volatility — not just mean draw — matters for grid-facing operation and demand response. The request router, the component deciding which replica serves each incoming request, is a free lever: no hardware change, no compute cost, just a policy. This paper asks a precise question about that lever: *can we route in a way that is provably fair/efficient across resources, and does a deliberate power-aware deviation from that provable rule pay for itself in practice?*

Contributions.

- (1) A formalization of power-aware LLM-serving routing as a 3-resource (compute/load/power) egalitarian social welfare problem, building on DRF [2] but extending it to a reactive, time-varying resource (live power-ramp state) rather than DRF’s original static per-request demand model — a real modeling departure, made explicit rather than assumed away (§3).
- (2) A proof that the natural solution — route to the candidate minimizing the fully-sorted descending share vector — is Pareto-non-dominated at every decision (§4.1).
- (3) A proof, via explicit counterexample, that a natural power-prioritized variant of this rule sacrifices the guarantee (§4.2) — and an argument for why the trade is nonetheless worth taking.
- (4) Hardware validation on a real 8×4090 LLM-serving fleet: the power-prioritized rule delivers a large, tightly-replicated win on tail power-ramp metrics under sustained pressure, at no measurable mean-latency cost (§5).

2 Background and Related Work

- **Dominant Resource Fairness (DRF)** [2] — multi-resource fair allocation via lexicographic comparison of dominant shares; proven to satisfy sharing incentive, envy-freeness, and Pareto efficiency for a fixed set of resources with static per-user demand vectors. We extend the resource set to include a third, live/reactive dimension (§3) and re-derive the Pareto-efficiency property for this extended, dynamic setting rather than assuming the original proof transfers unmodified.
- **LMETRIC** [4] — multiplicative $P\text{-token} \times BS$ scoring for LLM-serving load balancing; a simpler, non-fairness-theoretic alternative, useful as a baseline.
- **Power-of-Two-Choices** [1, 3] — sampled load balancing with a proven exponential improvement in expected max load; a different mechanism family (randomized sampling vs. our full-visibility deterministic rule), noted for completeness.
- **Coincidence factor / cross-machine power correlation** — the cross-GPU ramp-coincidence metric used in §5 is

drawn from a sibling paper’s grid-instrumentation framing, reused here as one of the two metrics the power-prioritized rule is evaluated on.

3 Problem Formulation

Consider N replicas serving requests behind a router. For a candidate replica c and an incoming request, define three shares, each normalized to roughly $[0, \infty)$ with 1.0 representing “at capacity”:

- $Share_{compute}(c) = P\text{-token}(c)/token_budget(c)$ — P -token is the prefix-cache-discounted count of new prefill tokens this request would cost replica c (computed from a per-replica cache-state mirror), so this share already reflects cache locality, not raw prompt length.
- $Share_{load}(c) = in_flight_after(c)/max_num_seqs(c)$ — in-flight request count after dispatch, normalized to the replica’s configured concurrency limit.
- $Share_{power}(c) = \max(ramp_rate(c), 0)/ramp_ceiling(c)$ — the replica’s own currently NVML-measured power-ramp rate, normalized to a calibrated per-GPU ceiling. Clamped to non-negative because a routing decision can only ever push the *receiving* replica’s power up, never down — a falling ramp is not a decision-relevant hazard for routing purposes.

Departure from DRF’s original model. Ghodsi et al.’s DRF [2] assumes a fixed, known demand vector per user, allocated from a fixed resource pool. $Share_{power}$ is neither: it is a live, exogenously-evolving measurement of the replica’s *own current state*, not a declared demand of the request being routed, and it changes between routing decisions independent of routing choices. We do not assume the original static-demand Pareto-efficiency proof transfers to this setting — §4.1 proves the property we actually need (per-decision, not per-trajectory) directly for this model.

Welfare objective. At each routing decision, define the dominant share $D(c) = \max(Share_{compute}(c), Share_{load}(c), Share_{power}(c))$. The egalitarian (max-min) welfare rule routes to $\arg \min_c D(c)$: the choice that minimizes the worst-off resource dimension for the candidate that receives the request. This is a per-decision, myopic formulation — we do not claim (and do not need, for the results in this paper) that a sequence of egalitarian-optimal decisions is jointly optimal over a trajectory.

4 Theory

4.1 The sorted-lexicographic rule is Pareto-non-dominated

$\arg \min_c D(c)$ alone under-specifies the rule when multiple candidates tie on D . Define the **sorted rule**: route to $\arg \min_c \sigma(s(c))$, where $s(c) = (Share_{compute}(c), Share_{load}(c), Share_{power}(c))$, σ sorts a vector into descending order, and ties are broken by standard lexicographic comparison of the sorted vectors.

LEMMA 4.1. *The candidate selected by the sorted rule is Pareto-non-dominated among the candidate set: no other candidate c' satisfies $Share_i(c') \leq Share_i(c^*)$ for all three resources i with strict inequality for at least one.*

PROOF. Suppose, for contradiction, some c' dominates the selected c^* . Then $s(c') \leq s(c^*)$ coordinatewise with strict inequality somewhere. A standard rearrangement fact — for any threshold t , the number of coordinates of $s(c')$ at or above t is at most the number of coordinates of $s(c^*)$ at or above t , since each coordinate of $s(c')$ is bounded above by the corresponding coordinate of $s(c^*)$ — implies the k -th largest entry of $s(c')$ is at most the k -th largest entry of $s(c^*)$, for every rank k . So $\sigma(s(c'))$ is coordinatewise $\leq \sigma(s(c^*))$. Summing coordinates (sort preserves sum) and using the strict inequality in the original domination, $\sum s(c') < \sum s(c^*)$ strictly, so $\sigma(s(c'))$ and $\sigma(s(c^*))$ cannot be equal as vectors. Combined with the coordinatewise $\leq$ just shown, the first rank at which they differ must favor c' strictly, so $\sigma(s(c')) <_{\text{lex}} \sigma(s(c^*))$. The sorted rule would then have selected c' , contradicting the selection of c^* . $\square$

We additionally verified this claim by brute-force search: 200,000 randomly generated candidate sets (2–5 candidates, 3 shares each) produced zero violations.

4.2 A power-prioritized variant sacrifices the guarantee

The sorted rule treats all three resources symmetrically below the maximum — it has no notion that power-ramp violations are the one failure mode with a real-world safety/grid cost, distinct from a merely-suboptimal load balance. A natural fix: after comparing the dominant share, always compare $Share_{power}$ next, *by name*, ahead of $Share_{load}$, regardless of which is numerically larger. Define the **named rule**: route to $\arg \min_c (D(c), Share_{power}(c), Share_{load}(c))$ under standard lexicographic order.

The named rule can select a Pareto-dominated candidate.

Construction. Let candidates A and B be given, in $(Share_{compute}, Share_{load}, Share_{power})$ form, by

$$A = (0.3, 0.9, 0.9), \quad B = (0.5, 0.9, 0.9).$$

A Pareto-dominates B (equal load and power, strictly lower compute). Both have $D(A) = D(B) = 0.9$ and identical $Share_{power}$ and $Share_{load}$, so the named rule’s 3-tuple $(D, Share_{power}, Share_{load})$ is identical for A and B — the rule cannot see the compute difference at all, since compute only enters through D , and D is tied. $\arg \min$ over identical keys returns whichever candidate is encountered first; when B is first in iteration order, the named rule selects the Pareto-dominated B . (Verified directly in code, not just argued: see `verify_pareto_lemma.py` in the project research log.)

Why this happens, and why the trade can still be worth it. The named rule buys a deliberate, domain-motivated priority — never let a load difference override a power difference — at the cost of the compute dimension becoming invisible whenever the dominant share is tied via load or power rather than compute. This is not a bug to be patched away casually: fixing it (e.g. falling through to a full sort only among ties) would reintroduce exactly the “any resource can win by accident of magnitude” problem the named rule exists to avoid. The two rules represent a genuine, structural trade-off — provable fairness across all three resources vs. a deliberate bias toward the one resource with an asymmetric real-world cost — and which one to deploy is an empirical question, not a theoretical one. §5 answers it.

Table 1: Sorted rule (Lemma-1-safe, drf_fixed) vs. named rule (power-prioritized, drf_power_tiebreak), mean $\pm$ std over 3 replicated trials.

Metric	Sorted	Named	Direction
Max ramp (W/s)	5010.6 $\pm$ 1940.8	3015.3 $\pm$ 257.1	−40%, 8 $\times$ tighter
Coincidence (%)	7.7 $\pm$ 1.5	4.4 $\pm$ 1.2	−43%
TBT max (ms)	2703.7 $\pm$ 968.5	2034.6 $\pm$ 181.6	−25%, more stable
Duty cycle	0.281 $\pm$ 0.021	0.259 $\pm$ 0.010	improved
TTFT mean (s)	3.712 $\pm$ 0.935	3.942 $\pm$ 0.271	flat, tighter std
Mean ramp (W/s)	191.86 $\pm$ 9.52	204.13 $\pm$ 7.51	+6%, minor cost

5 Experimental Validation

Setup. 8 $\times$ 4090 server (single chassis), 7B model (Qwen2.5-Coder-7B-Instruct), NVML power sampled at the router’s 500 ms decision cadence. Workload: whale-injection (15% long-prompt fraction, 13.6–15.5k-token whales, single-turn), request rate matched to real BurstGPT aggregate token throughput (not raw request rate, which saturates the fleet and produces uninformative results) — the condition under which the fleet sustains real power pressure (duty cycle $\approx$ 26–28% of decisions occur while some replica is above its ramp ceiling), which is precisely the regime §4.2’s trade-off is about. 3 replicated trials, fixed seed.

Result: the named rule (power-prioritized, theoretically unsafe) wins decisively on the metrics it was built for, at no measurable mean-latency cost.

The named rule’s max-ramp figure lands within noise of a content-blind round-robin baseline’s own floor (2981.3 $\pm$ 255.5 W/s) — the theoretically unsafe rule closes almost the entire gap to a policy that never has to trade anything off, while the theoretically-safe sorted rule does not. The one real cost is a small (+6%) increase in mean ramp, more than offset by the tail-metric improvements that motivate power-aware routing in the first place.

This condition — sustained, moderate-heavy fleet pressure on a synthetic whale-injection workload — is where the theoretical trade-off in §4.2 is directly exercised. Other tested conditions, including light-load traffic where the fleet is never pressured and real BurstGPT traffic at various rates, are outside this paper’s current scope; see the project research log for that data if useful context is wanted.

6 Discussion and Limitations

- **Small- N / shared-PDU caveat:** 8 $\times$ 4090 in one chassis likely shares upstream PDU/PSU — not independent grid circuits. Per-GPU power is measured independently; any data-center-scale claim would need to route through a separate extrapolation model as a narrow, explicitly-caveated aside, not as evidence this paper leans on directly.
- **Scope of validation:** the experimental result in §5 is demonstrated under sustained fleet power pressure on a controlled synthetic workload — the condition that directly exercises the theoretical trade-off characterized in §4.2. Generalization to arbitrary real-world traffic patterns is not claimed here and is a direction for future work.
- **Per-decision vs. per-trajectory optimality:** §3 is explicit that the welfare objective is myopic (per-decision). We do not claim, and §4.1 does not require, that a sequence of such decisions is optimal in aggregate over a trajectory

— only that each individual decision satisfies (or, for the named rule, deliberately trades away) a well-defined static guarantee.

7 Conclusion

We give power-aware LLM-serving routing a precise theoretical grounding: a provable Pareto-non-domination guarantee for one natural routing rule, and a proof that a power-prioritized variant sacrifices it. Rather than treating the sacrifice as a flaw, we show it is the right trade to make: on real 8 $\times$ 4090 hardware under sustained power pressure, the theoretically “unsafe” rule delivers large, tightly-replicated improvements on the tail power-ramp metrics that motivate power-aware routing in the first place, at no measurable cost to mean latency. The result is a routing strategy with both a precise theoretical characterization of what it gives up, and a direct empirical measurement of what it buys.

References

- [1] Yossi Azar, Andrei Z. Broder, Anna M. Karlin, and Eli Upfal. 1994. Balanced Allocations. In *Proceedings of the 26th Annual ACM Symposium on Theory of Computing (STOC '94)*. ACM.
- [2] Ali Ghodsi, Matei Zaharia, Benjamin Hindman, Andy Konwinski, Scott Shenker, and Ion Stoica. 2011. Dominant Resource Fairness: Fair Allocation of Multiple Resource Types. In *Proceedings of the 8th USENIX Symposium on Networked Systems Design and Implementation (NSDI '11)*. USENIX Association.
- [3] Michael Mitzenmacher. 2001. The Power of Two Choices in Randomized Load Balancing. *IEEE Transactions on Parallel and Distributed Systems* 12, 10 (2001), 1094–1104.
- [4] Dingyan Zhang, Jinbo Han, Kaixi Zhang, Xingda Wei, Sijie Shen, Chenguang Fang, Wenyuan Yu, Jingren Zhou, and Rong Chen. 2026. LMetric: Simple is Better — Multiplication May Be All You Need for LLM Request Scheduling. In *Proceedings of the 20th USENIX Symposium on Operating Systems Design and Implementation (OSDI '26)*. USENIX Association. arXiv:2603.15202.